

# 15\_CNN\_MNIST

December 1, 2020

## 1 3,5,7 Layers CNN on MNIST dataset

```
In [1]: from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Activation
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
```

```
# the data, split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
print("Number of X_train points :", X_train.shape)
print("Number of y_train points :", y_train.shape)
print("Number of X_test points :", X_test.shape)
print("Number of y_test points :", y_test.shape)
```

Number of X\_train points : (60000, 28, 28)

Number of y\_train points : (60000,)

Number of X\_test points : (10000, 28, 28)

Number of y\_test points : (10000,)

```
In [3]: #Initialization
```

```
batch_size = 128
```

```
num_classes = 10
```

```
epochs = 15
```

```
# input image dimensions
```

```
img_rows, img_cols = 28, 28
```

```
if K.image_data_format() == 'channels_first':
```

```
    X_train = X_train.reshape(X_train.shape[0], 1, img_rows, img_cols)
```

```
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
```

```
    input_shape = (1, img_rows, img_cols)
```

```
else:
```

```
    X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1) # 1 here is even
```

```

X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)

X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255
X_test /= 255
print('X_train shape:', X_train.shape)
print(X_train.shape[0], 'train samples')
print(X_test.shape[0], 'test samples')

X_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

In [4]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = keras.utils.to_categorical(y_train, 10)
Y_test = keras.utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```

## 1.1 3 layer CNN

```

In [5]: from keras.layers.normalization import BatchNormalization
        #Activation=ReLU
        #Optimizer=Adam
        #Architecture= conv2d,maxpool,conv2d,maxpool,conv2d,bn,maxpool,fc,dropout,sm

model_3l = Sequential()
# first layer
model_3l.add(Conv2D(32, kernel_size=(3, 3),activation='relu',input_shape=input_shape))
model_3l.add(MaxPooling2D(pool_size=(2, 2)))
# second layer
model_3l.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model_3l.add(MaxPooling2D(pool_size=(2, 2)))
# third layer
model_3l.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model_3l.add(BatchNormalization(axis=1))

```

```

model_3l.add(MaxPooling2D(pool_size=(2, 2)))
# flattening, dropout
model_3l.add(Flatten())
model_3l.add(Dense(128, activation='relu'))
model_3l.add(Dropout(0.5))
# softmax
model_3l.add(Dense(num_classes, activation='softmax'))
model_3l.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
batch_normalization (Batch Normalization)	(None, 3, 3, 64)	12
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 64)	0
flatten (Flatten)	(None, 64)	0
dense (Dense)	(None, 128)	8320
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
Total params: 65,366		
Trainable params: 65,360		
Non-trainable params: 6		

```

In [6]: model_3l.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam,
        history=model_3l.fit(X_train, Y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(X_val, Y_val)))

```

Epoch 1/15

469/469 [=====] - 3s 6ms/step - loss: 0.3624 - accuracy: 0.8861 - val\_loss: 0.1007

Epoch 2/15

469/469 [=====] - 2s 5ms/step - loss: 0.0907 - accuracy: 0.9735 - val\_loss: 0.0100

```

Epoch 3/15
469/469 [=====] - 2s 5ms/step - loss: 0.0657 - accuracy: 0.9811 - val.
Epoch 4/15
469/469 [=====] - 2s 5ms/step - loss: 0.0529 - accuracy: 0.9848 - val.
Epoch 5/15
469/469 [=====] - 2s 5ms/step - loss: 0.0452 - accuracy: 0.9861 - val.
Epoch 6/15
469/469 [=====] - 2s 5ms/step - loss: 0.0367 - accuracy: 0.9892 - val.
Epoch 7/15
469/469 [=====] - 2s 5ms/step - loss: 0.0339 - accuracy: 0.9896 - val.
Epoch 8/15
469/469 [=====] - 2s 5ms/step - loss: 0.0283 - accuracy: 0.9916 - val.
Epoch 9/15
469/469 [=====] - 2s 5ms/step - loss: 0.0255 - accuracy: 0.9918 - val.
Epoch 10/15
469/469 [=====] - 2s 5ms/step - loss: 0.0225 - accuracy: 0.9930 - val.
Epoch 11/15
469/469 [=====] - 2s 5ms/step - loss: 0.0224 - accuracy: 0.9931 - val.
Epoch 12/15
469/469 [=====] - 2s 5ms/step - loss: 0.0174 - accuracy: 0.9944 - val.
Epoch 13/15
469/469 [=====] - 2s 5ms/step - loss: 0.0164 - accuracy: 0.9949 - val.
Epoch 14/15
469/469 [=====] - 2s 5ms/step - loss: 0.0160 - accuracy: 0.9949 - val.
Epoch 15/15
469/469 [=====] - 2s 5ms/step - loss: 0.0142 - accuracy: 0.9954 - val.

```

```

In [7]: #plotting function
        %matplotlib notebook
        import matplotlib.pyplot as plt
        %matplotlib inline
        import numpy as np

        def plt_dynamic(x, vy, ty, ax, colors=['b']):
            ax.plot(x, vy, 'b', label="Validation Loss")
            ax.plot(x, ty, 'r', label="Train Loss")
            plt.legend()
            plt.grid()
            fig.canvas.draw()

In [8]: score = model_3l.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])
        print('Test accuracy:', score[1])

        fig,ax = plt.subplots(1,1)
        ax.set_title('EpochsVS Loss')

```

```

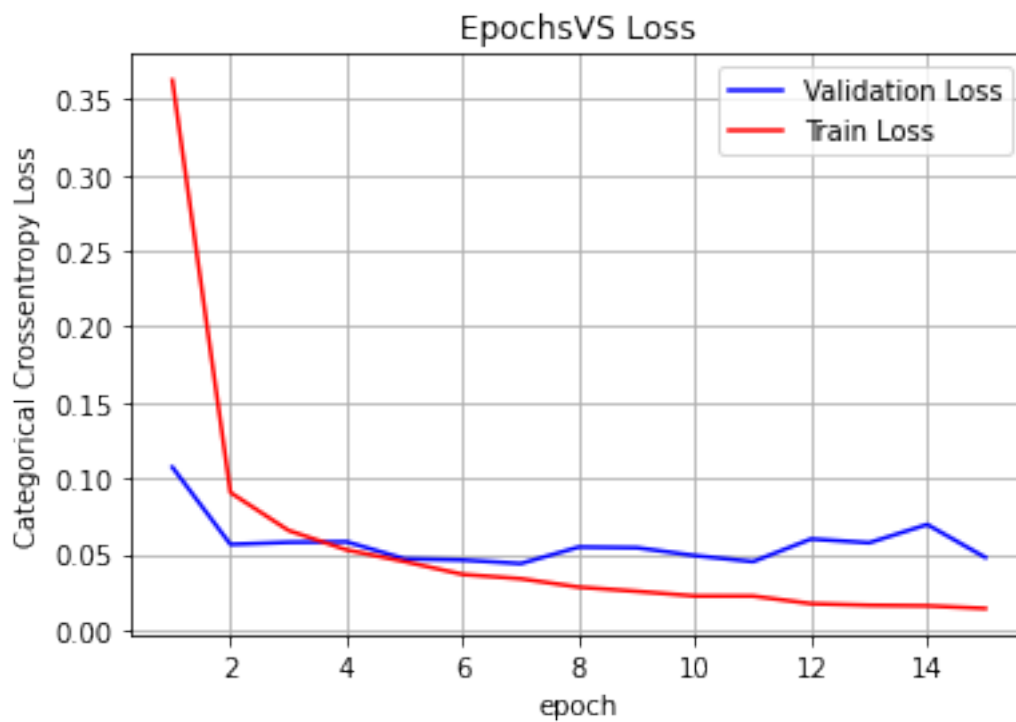
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.047884970903396606  
Test accuracy: 0.9900000095367432



## 1.2 5 layer CNN

```

In [9]: #Architecture = Conv2d,pool,conv2d,d,pool,conv2d,bn,pool,conv2d,bn,pool,conv2d,pool
model_5l = Sequential()
# first layer
model_5l.add(Conv2D(64, kernel_size=(2, 2),activation='relu',input_shape=input_shape))
model_5l.add(MaxPooling2D(pool_size=(1, 1)))
# second layer
model_5l.add(Conv2D(64, kernel_size=(2, 2), activation='relu'))
model_5l.add(Dropout(0.25))

```

```

model_51.add(MaxPooling2D(pool_size=(2, 2)))
# third layer
model_51.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model_51.add(BatchNormalization(axis=1))
model_51.add(MaxPooling2D(pool_size=(1, 1)))
# fourth layer
model_51.add(Conv2D(32, kernel_size=(2, 2), activation='relu'))
model_51.add(BatchNormalization(axis=1))
model_51.add(MaxPooling2D(pool_size=(2, 2)))
# fifth layer
model_51.add(Conv2D(16, kernel_size=(2, 2), activation='relu'))
model_51.add(MaxPooling2D(pool_size=(2, 2)))
# flattening, dropout
model_51.add(Flatten())
model_51.add(Dense(128, activation='relu'))
model_51.add(Dropout(0.5))
# softmax
model_51.add(Dense(num_classes, activation='softmax'))
model_51.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 27, 27, 64)	320
max_pooling2d_3 (MaxPooling2D)	(None, 27, 27, 64)	0
conv2d_4 (Conv2D)	(None, 26, 26, 64)	16448
dropout_1 (Dropout)	(None, 26, 26, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 13, 13, 64)	0
conv2d_5 (Conv2D)	(None, 11, 11, 32)	18464
batch_normalization_1 (Batch Normalization)	(None, 11, 11, 32)	44
max_pooling2d_5 (MaxPooling2D)	(None, 11, 11, 32)	0
conv2d_6 (Conv2D)	(None, 10, 10, 32)	4128
batch_normalization_2 (Batch Normalization)	(None, 10, 10, 32)	40
max_pooling2d_6 (MaxPooling2D)	(None, 5, 5, 32)	0
conv2d_7 (Conv2D)	(None, 4, 4, 16)	2064

max_pooling2d_7 (MaxPooling2)	(None, 2, 2, 16)	0
-----		
flatten_1 (Flatten)	(None, 64)	0
-----		
dense_2 (Dense)	(None, 128)	8320
-----		
dropout_2 (Dropout)	(None, 128)	0
-----		
dense_3 (Dense)	(None, 10)	1290
=====		
Total params: 51,118		
Trainable params: 51,076		
Non-trainable params: 42		
-----		

```
In [10]: model_51.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam,
        history=model_51.fit(X_train, Y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(X_val, Y_val)))
```

```
Epoch 1/15
469/469 [=====] - 5s 11ms/step - loss: 0.4308 - accuracy: 0.8625 - val_loss: 0.4308 - val_accuracy: 0.8625
Epoch 2/15
469/469 [=====] - 5s 11ms/step - loss: 0.1072 - accuracy: 0.9684 - val_loss: 0.1072 - val_accuracy: 0.9684
Epoch 3/15
469/469 [=====] - 5s 11ms/step - loss: 0.0763 - accuracy: 0.9771 - val_loss: 0.0763 - val_accuracy: 0.9771
Epoch 4/15
469/469 [=====] - 5s 11ms/step - loss: 0.0623 - accuracy: 0.9813 - val_loss: 0.0623 - val_accuracy: 0.9813
Epoch 5/15
469/469 [=====] - 5s 11ms/step - loss: 0.0537 - accuracy: 0.9845 - val_loss: 0.0537 - val_accuracy: 0.9845
Epoch 6/15
469/469 [=====] - 5s 11ms/step - loss: 0.0483 - accuracy: 0.9857 - val_loss: 0.0483 - val_accuracy: 0.9857
Epoch 7/15
469/469 [=====] - 5s 11ms/step - loss: 0.0427 - accuracy: 0.9871 - val_loss: 0.0427 - val_accuracy: 0.9871
Epoch 8/15
469/469 [=====] - 5s 11ms/step - loss: 0.0396 - accuracy: 0.9873 - val_loss: 0.0396 - val_accuracy: 0.9873
Epoch 9/15
469/469 [=====] - 5s 11ms/step - loss: 0.0376 - accuracy: 0.9887 - val_loss: 0.0376 - val_accuracy: 0.9887
Epoch 10/15
469/469 [=====] - 5s 11ms/step - loss: 0.0330 - accuracy: 0.9900 - val_loss: 0.0330 - val_accuracy: 0.9900
Epoch 11/15
469/469 [=====] - 5s 11ms/step - loss: 0.0323 - accuracy: 0.9904 - val_loss: 0.0323 - val_accuracy: 0.9904
Epoch 12/15
469/469 [=====] - 5s 11ms/step - loss: 0.0303 - accuracy: 0.9909 - val_loss: 0.0303 - val_accuracy: 0.9909
Epoch 13/15
469/469 [=====] - 5s 11ms/step - loss: 0.0280 - accuracy: 0.9912 - val_loss: 0.0280 - val_accuracy: 0.9912
Epoch 14/15
469/469 [=====] - 5s 11ms/step - loss: 0.0272 - accuracy: 0.9919 - val_loss: 0.0272 - val_accuracy: 0.9919
Epoch 15/15
```

469/469 [=====] - 5s 11ms/step - loss: 0.0235 - accuracy: 0.9923 - va

```
In [11]: score = model_51.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```


```
fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.026484984904527664

Test accuracy: 0.9919999837875366





### 1.3 7 layer CNN

```
In [12]: #Architecture=conv2d,pool,conv2d,bn,pool,conv2d,d,pool,conv2d,d,pool,conv2d,bn,pool,c
model_71 = Sequential()
# first layer
model_71.add(Conv2D(64, kernel_size=(5, 5),activation='relu',input_shape=input_shape))
model_71.add(MaxPooling2D(pool_size=(1, 1)))
# second layer
model_71.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model_71.add(BatchNormalization(axis=1))
model_71.add(MaxPooling2D(pool_size=(2, 2)))
# third layer
model_71.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model_71.add(Dropout(0.5))
model_71.add(MaxPooling2D(pool_size=(1, 1)))
# fourth layer
model_71.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model_71.add(Dropout(0.25))
model_71.add(MaxPooling2D(pool_size=(1, 1)))
# fifth layer
model_71.add(Conv2D(16, kernel_size=(2, 2), activation='relu'))
model_71.add(BatchNormalization(axis=1))
model_71.add(MaxPooling2D(pool_size=(1, 1)))
# sixth layer
model_71.add(Conv2D(16, kernel_size=(2, 2), activation='relu'))
model_71.add(BatchNormalization(axis=1))
model_71.add(MaxPooling2D(pool_size=(1, 1)))
# seventh layer
model_71.add(Conv2D(8, kernel_size=(3, 3), activation='relu'))
model_71.add(MaxPooling2D(pool_size=(1, 1)))
# flattening,dropout
model_71.add(Flatten())
model_71.add(Dense(128, activation='relu'))
model_71.add(Dropout(0.5))
# softmax
model_71.add(Dense(num_classes, activation='softmax'))
model_71.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 24, 24, 64)	1664
max_pooling2d_8 (MaxPooling2D)	(None, 24, 24, 64)	0
conv2d_9 (Conv2D)	(None, 22, 22, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 22, 22, 64)	88

max_pooling2d_9	(MaxPooling2 (None, 11, 11, 64)	0
conv2d_10	(Conv2D) (None, 9, 9, 32)	18464
dropout_3	(Dropout) (None, 9, 9, 32)	0
max_pooling2d_10	(MaxPooling (None, 9, 9, 32)	0
conv2d_11	(Conv2D) (None, 7, 7, 32)	9248
dropout_4	(Dropout) (None, 7, 7, 32)	0
max_pooling2d_11	(MaxPooling (None, 7, 7, 32)	0
conv2d_12	(Conv2D) (None, 6, 6, 16)	2064
batch_normalization_4	(Batch (None, 6, 6, 16)	24
max_pooling2d_12	(MaxPooling (None, 6, 6, 16)	0
conv2d_13	(Conv2D) (None, 5, 5, 16)	1040
batch_normalization_5	(Batch (None, 5, 5, 16)	20
max_pooling2d_13	(MaxPooling (None, 5, 5, 16)	0
conv2d_14	(Conv2D) (None, 3, 3, 8)	1160
max_pooling2d_14	(MaxPooling (None, 3, 3, 8)	0
flatten_2	(Flatten) (None, 72)	0
dense_4	(Dense) (None, 128)	9344
dropout_5	(Dropout) (None, 128)	0
dense_5	(Dense) (None, 10)	1290
=====		
Total params: 81,334		
Trainable params: 81,268		
Non-trainable params: 66		
-----		

```
In [13]: model_71.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam,
          history=model_71.fit(X_train, Y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(X_val, Y_val)))

Epoch 1/15
```

```

469/469 [=====] - 6s 12ms/step - loss: 0.5877 - accuracy: 0.8027 - va
Epoch 2/15
469/469 [=====] - 5s 12ms/step - loss: 0.1203 - accuracy: 0.9651 - va
Epoch 3/15
469/469 [=====] - 5s 12ms/step - loss: 0.0888 - accuracy: 0.9741 - va
Epoch 4/15
469/469 [=====] - 6s 12ms/step - loss: 0.0748 - accuracy: 0.9788 - va
Epoch 5/15
469/469 [=====] - 5s 12ms/step - loss: 0.0676 - accuracy: 0.9810 - va
Epoch 6/15
469/469 [=====] - 6s 12ms/step - loss: 0.0589 - accuracy: 0.9826 - va
Epoch 7/15
469/469 [=====] - 5s 12ms/step - loss: 0.0583 - accuracy: 0.9834 - va
Epoch 8/15
469/469 [=====] - 5s 12ms/step - loss: 0.0532 - accuracy: 0.9855 - va
Epoch 9/15
469/469 [=====] - 5s 12ms/step - loss: 0.0496 - accuracy: 0.9857 - va
Epoch 10/15
469/469 [=====] - 6s 12ms/step - loss: 0.0451 - accuracy: 0.9868 - va
Epoch 11/15
469/469 [=====] - 5s 12ms/step - loss: 0.0434 - accuracy: 0.9869 - va
Epoch 12/15
469/469 [=====] - 5s 12ms/step - loss: 0.0416 - accuracy: 0.9886 - va
Epoch 13/15
469/469 [=====] - 5s 12ms/step - loss: 0.0386 - accuracy: 0.9889 - va
Epoch 14/15
469/469 [=====] - 5s 12ms/step - loss: 0.0374 - accuracy: 0.9891 - va
Epoch 15/15
469/469 [=====] - 5s 12ms/step - loss: 0.0372 - accuracy: 0.9891 - va

```

```

In [14]: score = model_71.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

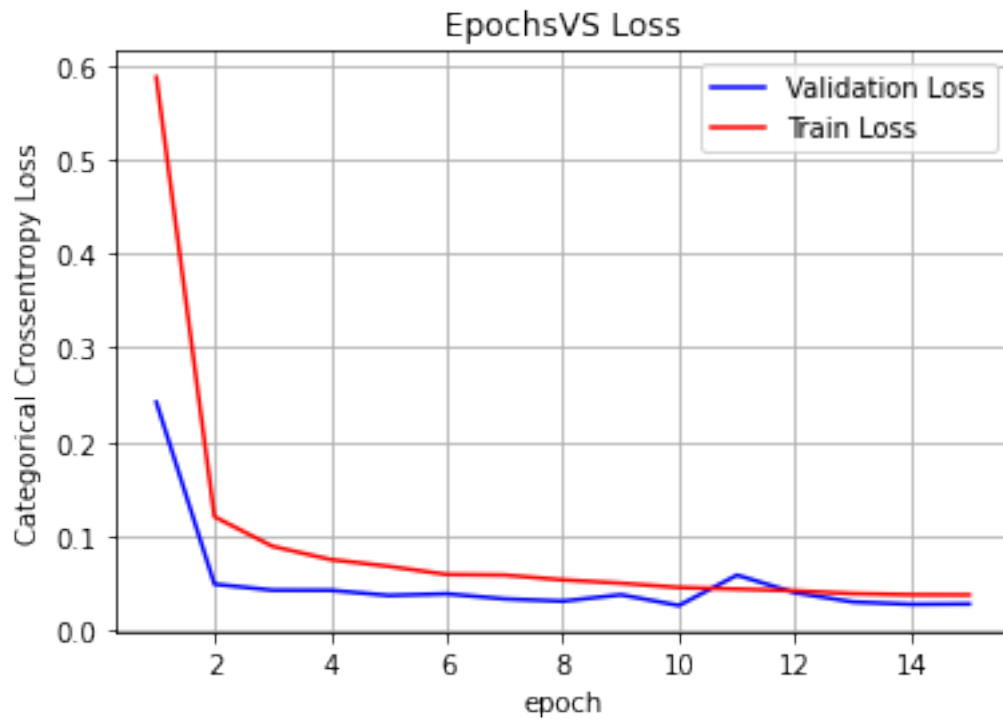
# list of epoch numbers
x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.027606463059782982

Test accuracy: 0.9926999807357788



```
In [16]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["No. of layer", "Accuracy %"]

x.add_row(["3", "99.00"])
x.add_row(["5", "99.19"])
x.add_row(["7", "99.26"])
print(x)
```

```
+-----+-----+
| No. of layer | Accuracy % |
+-----+-----+
|      3      |    99.00    |
|      5      |    99.19    |
|      7      |    99.26    |
+-----+-----+
```

## 2 Summary

- 1 created 3,5,7 layered cnn using batch normalization
- 2 observed that as layer increases accuracy also increased