

14_keras_mnist

November 30, 2020

0.1 2,3,5 hidden layer architecture on MNIST dataset

```
In [1]: import tensorflow as tf
        from tensorflow.keras import utils
        from tensorflow.keras.datasets import mnist
        import seaborn as sns
        from tensorflow.keras.initializers import he_normal
```

```
In [2]: # the data, shuffled and split between train and test sets
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

```
In [3]: print("Number of X_train points :", X_train.shape)
        print("Number of y_train points :", y_train.shape)
        print("Number of X_test points :", X_test.shape)
        print("Number of y_test points :", y_test.shape)
```

```
Number of X_train points : (60000, 28, 28)
Number of y_train points : (60000,)
Number of X_test points : (10000, 28, 28)
Number of y_test points : (10000,)
```

```
In [4]: # if you observe the input shape its 2 dimensional vector
        # for each image we have a (28*28) vector
        # we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [5]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (",
print("Number of training examples :", X_test.shape[0], "and each image is of shape (",
```

```
Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)
```

```
In [6]: # An example data point
        print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
   0  0  0  0  0  49 238 253 253 253 253 253 253 253 253 251  93  82
  82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0 16  93 252 253 187
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
  80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
```

```
In [7]: # if we observe the above matrix each cell is having a value between 0-255
        # before we move to apply machine learning algorithms lets try to normalize the data
        #  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

```
In [8]: # example data point after normalizing
print(X_train[0])
```

[illegible]

0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]

In [9]: # here we are having a class number for each image

```

print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = utils.to_categorical(y_train, 10)
Y_test = utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

```

```

Class label of first image : 5
After converting the output into a vector :  [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```

```

In [10]: # some model parameters
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

0.2 Two hidden layer

0.2.1 MLP + ReLU + ADAM + BN + w/o Dropout +2Layer

```

In [11]: from tensorflow.keras.layers import BatchNormalization
         initializer = tf.keras.initializers.RandomNormal(mean=0.0, stddev=0.066, seed=None)

         model_2lbn = Sequential() #2lbn=2 layer batch normalization

         model_2lbn.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=initializer))
         model_2lbn.add(BatchNormalization())

         model_2lbn.add(Dense(128, activation='relu', kernel_initializer=initializer))
         model_2lbn.add(BatchNormalization())

         model_2lbn.add(Dense(output_dim, activation='softmax'))

         model_2lbn.summary()

```

```
Model: "sequential"
```

```

-----
Layer (type)                Output Shape                Param #
=====

```

dense (Dense)	(None, 512)	401920

batch_normalization (Batch Normalization)	(None, 512)	2048

dense_1 (Dense)	(None, 128)	65664

batch_normalization_1 (Batch Normalization)	(None, 128)	512

dense_2 (Dense)	(None, 10)	1290
=====		
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

```
In [12]: model_2lbn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_2lbn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
Epoch 1/20
469/469 [=====] - 4s 9ms/step - loss: 0.1957 - accuracy: 0.9419 - val_loss: 0.1957
Epoch 2/20
469/469 [=====] - 4s 9ms/step - loss: 0.0712 - accuracy: 0.9786 - val_loss: 0.0712
Epoch 3/20
469/469 [=====] - 4s 9ms/step - loss: 0.0444 - accuracy: 0.9864 - val_loss: 0.0444
Epoch 4/20
469/469 [=====] - 4s 9ms/step - loss: 0.0308 - accuracy: 0.9905 - val_loss: 0.0308
Epoch 5/20
469/469 [=====] - 4s 9ms/step - loss: 0.0220 - accuracy: 0.9932 - val_loss: 0.0220
Epoch 6/20
469/469 [=====] - 4s 9ms/step - loss: 0.0213 - accuracy: 0.9932 - val_loss: 0.0213
Epoch 7/20
469/469 [=====] - 4s 9ms/step - loss: 0.0201 - accuracy: 0.9932 - val_loss: 0.0201
Epoch 8/20
469/469 [=====] - 4s 9ms/step - loss: 0.0136 - accuracy: 0.9956 - val_loss: 0.0136
Epoch 9/20
469/469 [=====] - 4s 9ms/step - loss: 0.0156 - accuracy: 0.9947 - val_loss: 0.0156
Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0140 - accuracy: 0.9952 - val_loss: 0.0140
Epoch 11/20
469/469 [=====] - 4s 9ms/step - loss: 0.0090 - accuracy: 0.9973 - val_loss: 0.0090
Epoch 12/20
469/469 [=====] - 4s 9ms/step - loss: 0.0115 - accuracy: 0.9964 - val_loss: 0.0115
Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0123 - accuracy: 0.9961 - val_loss: 0.0123
Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0088 - accuracy: 0.9973 - val_loss: 0.0088
```

```

Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0084 - accuracy: 0.9973 - val.
Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0065 - accuracy: 0.9977 - val.
Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.0116 - accuracy: 0.9963 - val.
Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0094 - accuracy: 0.9970 - val.
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0064 - accuracy: 0.9978 - val.
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0073 - accuracy: 0.9975 - val.

```

```

In [13]: #plotting function
         %matplotlib notebook
         import matplotlib.pyplot as plt
         %matplotlib inline
         import numpy as np

         def plt_dynamic(x, vy, ty, ax, colors=['b']):
             ax.plot(x, vy, 'b', label="Validation Loss")
             ax.plot(x, ty, 'r', label="Train Loss")
             plt.legend()
             plt.grid()
             fig.canvas.draw()

In [14]: score = model_2lbn.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_title('EpochsVS Loss')
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))


         vy = history.history['val_loss']
         ty = history.history['loss']
         plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.08280453830957413
Test accuracy: 0.9803000092506409

```

0.2.2 MLP + ReLU + ADAM + BN + with Dropout + 2Layer

```
In [15]: from tensorflow.keras.layers import Dropout
         from tensorflow.keras.layers import BatchNormalization
         initializer = tf.keras.initializers.he_normal(seed=None)

         model_2lbnd = Sequential() #model_2lbnd=2 layer batch normalization dropout

         model_2lbnd.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=initializer))
         model_2lbnd.add(BatchNormalization())
         model_2lbnd.add(Dropout(0.5))

         model_2lbnd.add(Dense(128, activation='relu', kernel_initializer=initializer))
         model_2lbnd.add(BatchNormalization())
         model_2lbnd.add(Dropout(0.5))

         model_2lbnd.add(Dense(output_dim, activation='softmax'))

         model_2lbnd.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
dense_3 (Dense)                (None, 512)                401920
-----
batch_normalization_2 (Batch Normalization) (None, 512)                2048
-----
dropout (Dropout)              (None, 512)                 0
-----
dense_4 (Dense)                (None, 128)                65664
-----
batch_normalization_3 (Batch Normalization) (None, 128)                512
-----
dropout_1 (Dropout)            (None, 128)                 0
-----
dense_5 (Dense)                (None, 10)                 1290
=====
Total params: 471,434
Trainable params: 470,154
Non-trainable params: 1,280
-----

```

```
In [16]: model_2lbnd.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_2lbnd.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_val, Y_val))
```

```

Epoch 1/20
469/469 [=====] - 5s 10ms/step - loss: 0.4309 - accuracy: 0.8677 - val_loss: 0.3809 - val_accuracy: 0.8800
Epoch 2/20
469/469 [=====] - 5s 10ms/step - loss: 0.2066 - accuracy: 0.9382 - val_loss: 0.2066 - val_accuracy: 0.9382
Epoch 3/20
469/469 [=====] - 5s 10ms/step - loss: 0.1592 - accuracy: 0.9523 - val_loss: 0.1592 - val_accuracy: 0.9523
Epoch 4/20
469/469 [=====] - 5s 10ms/step - loss: 0.1346 - accuracy: 0.9592 - val_loss: 0.1346 - val_accuracy: 0.9592
Epoch 5/20
469/469 [=====] - 5s 10ms/step - loss: 0.1214 - accuracy: 0.9624 - val_loss: 0.1214 - val_accuracy: 0.9624
Epoch 6/20
469/469 [=====] - 5s 10ms/step - loss: 0.1101 - accuracy: 0.9665 - val_loss: 0.1101 - val_accuracy: 0.9665
Epoch 7/20
469/469 [=====] - 5s 10ms/step - loss: 0.1000 - accuracy: 0.9691 - val_loss: 0.1000 - val_accuracy: 0.9691
Epoch 8/20
469/469 [=====] - 5s 10ms/step - loss: 0.0965 - accuracy: 0.9702 - val_loss: 0.0965 - val_accuracy: 0.9702
Epoch 9/20
469/469 [=====] - 5s 10ms/step - loss: 0.0853 - accuracy: 0.9739 - val_loss: 0.0853 - val_accuracy: 0.9739
Epoch 10/20
469/469 [=====] - 5s 10ms/step - loss: 0.0831 - accuracy: 0.9740 - val_loss: 0.0831 - val_accuracy: 0.9740
Epoch 11/20
469/469 [=====] - 5s 10ms/step - loss: 0.0774 - accuracy: 0.9759 - val_loss: 0.0774 - val_accuracy: 0.9759
Epoch 12/20

```

```

469/469 [=====] - 5s 10ms/step - loss: 0.0754 - accuracy: 0.9763 - va
Epoch 13/20
469/469 [=====] - 5s 10ms/step - loss: 0.0697 - accuracy: 0.9776 - va
Epoch 14/20
469/469 [=====] - 5s 10ms/step - loss: 0.0695 - accuracy: 0.9785 - va
Epoch 15/20
469/469 [=====] - 5s 10ms/step - loss: 0.0665 - accuracy: 0.9792 - va
Epoch 16/20
469/469 [=====] - 5s 10ms/step - loss: 0.0628 - accuracy: 0.9807 - va
Epoch 17/20
469/469 [=====] - 5s 10ms/step - loss: 0.0636 - accuracy: 0.9801 - va
Epoch 18/20
469/469 [=====] - 5s 10ms/step - loss: 0.0578 - accuracy: 0.9810 - va
Epoch 19/20
469/469 [=====] - 5s 10ms/step - loss: 0.0558 - accuracy: 0.9821 - va
Epoch 20/20
469/469 [=====] - 5s 10ms/step - loss: 0.0542 - accuracy: 0.9831 - va

```

```

In [17]: score = model_2lbnd.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.05603070184588432
Test accuracy: 0.9840999841690063

```



0.3 Three hidden layer

0.3.1 MLP + ReLU + ADAM + w/o Dropout + BN + 3Layer

```
In [ ]: initializer = tf.keras.initializers.he_normal(seed=None)
```

```
model_3lbn = Sequential()
```

```
model_3lbn.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=initializer))
model_3lbn.add(BatchNormalization())
```

```
model_3lbn.add(Dense(128, activation='relu', kernel_initializer=initializer) ) #layer 2
model_3lbn.add(BatchNormalization())
```

```
model_3lbn.add(Dense(100, activation='relu', kernel_initializer=initializer) ) #layer 3
model_3lbn.add(BatchNormalization())
```

```
model_3lbn.add(Dense(output_dim, activation='softmax'))
```

```
model_3lbn.summary()
```

```
Model: "sequential_6"
```

```
-----
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_9 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_10 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 100)	12900
batch_normalization_11 (Batch Normalization)	(None, 100)	400
dense_17 (Dense)	(None, 10)	1010
Total params: 484,454		
Trainable params: 482,974		
Non-trainable params: 1,480		

```
In [ ]: model_3lbn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_3lbn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

```
Epoch 1/20
469/469 [=====] - 6s 13ms/step - loss: 0.2015 - accuracy: 0.9404 - val_loss: 0.1815 - val_accuracy: 0.9500
Epoch 2/20
469/469 [=====] - 6s 12ms/step - loss: 0.0771 - accuracy: 0.9771 - val_loss: 0.0655 - val_accuracy: 0.9889
Epoch 3/20
469/469 [=====] - 6s 12ms/step - loss: 0.0515 - accuracy: 0.9839 - val_loss: 0.0455 - val_accuracy: 0.9900
Epoch 4/20
469/469 [=====] - 5s 12ms/step - loss: 0.0359 - accuracy: 0.9886 - val_loss: 0.0355 - val_accuracy: 0.9900
Epoch 5/20
469/469 [=====] - 5s 12ms/step - loss: 0.0292 - accuracy: 0.9908 - val_loss: 0.0305 - val_accuracy: 0.9900
Epoch 6/20
469/469 [=====] - 6s 12ms/step - loss: 0.0296 - accuracy: 0.9899 - val_loss: 0.0305 - val_accuracy: 0.9900
Epoch 7/20
469/469 [=====] - 6s 12ms/step - loss: 0.0211 - accuracy: 0.9932 - val_loss: 0.0255 - val_accuracy: 0.9900
Epoch 8/20
469/469 [=====] - 6s 12ms/step - loss: 0.0184 - accuracy: 0.9937 - val_loss: 0.0255 - val_accuracy: 0.9900
Epoch 9/20
469/469 [=====] - 5s 11ms/step - loss: 0.0213 - accuracy: 0.9929 - val_loss: 0.0255 - val_accuracy: 0.9900
Epoch 10/20
469/469 [=====] - 5s 12ms/step - loss: 0.0150 - accuracy: 0.9950 - val_loss: 0.0255 - val_accuracy: 0.9900
Epoch 11/20
469/469 [=====] - 5s 12ms/step - loss: 0.0137 - accuracy: 0.9956 - val_loss: 0.0255 - val_accuracy: 0.9900
```

```

Epoch 12/20
469/469 [=====] - 5s 12ms/step - loss: 0.0162 - accuracy: 0.9944 - va
Epoch 13/20
469/469 [=====] - 5s 12ms/step - loss: 0.0132 - accuracy: 0.9957 - va
Epoch 14/20
469/469 [=====] - 6s 12ms/step - loss: 0.0112 - accuracy: 0.9960 - va
Epoch 15/20
469/469 [=====] - 6s 12ms/step - loss: 0.0091 - accuracy: 0.9970 - va
Epoch 16/20
469/469 [=====] - 6s 12ms/step - loss: 0.0129 - accuracy: 0.9957 - va
Epoch 17/20
469/469 [=====] - 6s 12ms/step - loss: 0.0111 - accuracy: 0.9963 - va
Epoch 18/20
469/469 [=====] - 6s 12ms/step - loss: 0.0086 - accuracy: 0.9973 - va
Epoch 19/20
469/469 [=====] - 6s 12ms/step - loss: 0.0076 - accuracy: 0.9976 - va
Epoch 20/20
469/469 [=====] - 6s 12ms/step - loss: 0.0093 - accuracy: 0.9970 - va

```

```

In [ ]: score = model_3lbn.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])
        print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.105586476624012
Test accuracy: 0.9764000177383423

```



0.3.2 MLP + ReLU + ADAM + with Dropout + BN + 3Layer

```
In [ ]: initializer = tf.keras.initializers.he_normal(seed=None)
```

```
model_3lbnd = Sequential()
```

```
model_3lbnd.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=initializer))
```

```
model_3lbnd.add(BatchNormalization())
```

```
model_3lbnd.add(Dropout(0.5))
```

```
model_3lbnd.add(Dense(128, activation='relu', kernel_initializer=initializer)) #layer 2
```

```
model_3lbnd.add(BatchNormalization())
```

```
model_3lbnd.add(Dropout(0.5))
```

```
model_3lbnd.add(Dense(100, activation='relu', kernel_initializer=initializer)) #layer 3
```

```
model_3lbnd.add(BatchNormalization())
```

```
model_3lbnd.add(Dropout(0.5))
```

```
model_3lbnd.add(Dense(output_dim, activation='softmax'))
```

```
model_3lbnd.summary()
```

```
Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 512)	401920
batch_normalization_12 (Batch Normalization)	(None, 512)	2048
dropout_2 (Dropout)	(None, 512)	0
dense_19 (Dense)	(None, 128)	65664
batch_normalization_13 (Batch Normalization)	(None, 128)	512
dropout_3 (Dropout)	(None, 128)	0
dense_20 (Dense)	(None, 100)	12900
batch_normalization_14 (Batch Normalization)	(None, 100)	400
dropout_4 (Dropout)	(None, 100)	0
dense_21 (Dense)	(None, 10)	1010
Total params: 484,454		
Trainable params: 482,974		
Non-trainable params: 1,480		

```
In [ ]: model_3lbn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_3lbn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Epoch 1/20

469/469 [=====] - 6s 14ms/step - loss: 0.6442 - accuracy: 0.8026 - val_loss: 0.2782 - val_accuracy: 0.9201

Epoch 2/20

469/469 [=====] - 6s 13ms/step - loss: 0.2782 - accuracy: 0.9201 - val_loss: 0.2113 - val_accuracy: 0.9397

Epoch 3/20

469/469 [=====] - 6s 13ms/step - loss: 0.2113 - accuracy: 0.9397 - val_loss: 0.1809 - val_accuracy: 0.9477

Epoch 4/20

469/469 [=====] - 6s 13ms/step - loss: 0.1809 - accuracy: 0.9477 - val_loss: 0.1580 - val_accuracy: 0.9547

Epoch 5/20

469/469 [=====] - 6s 13ms/step - loss: 0.1580 - accuracy: 0.9547 - val_loss: 0.1457 - val_accuracy: 0.9582

Epoch 6/20

469/469 [=====] - 6s 13ms/step - loss: 0.1457 - accuracy: 0.9582 - val_loss: 0.1275 - val_accuracy: 0.9631

Epoch 7/20

469/469 [=====] - 6s 13ms/step - loss: 0.1275 - accuracy: 0.9631 - val_loss: 0.1150 - val_accuracy: 0.9700

Epoch 8/20


```

469/469 [=====] - 7s 14ms/step - loss: 0.1243 - accuracy: 0.9639 - va
Epoch 9/20
469/469 [=====] - 7s 14ms/step - loss: 0.1144 - accuracy: 0.9668 - va
Epoch 10/20
469/469 [=====] - 7s 16ms/step - loss: 0.1059 - accuracy: 0.9696 - va
Epoch 11/20
469/469 [=====] - 7s 14ms/step - loss: 0.1014 - accuracy: 0.9706 - va
Epoch 12/20
469/469 [=====] - 6s 13ms/step - loss: 0.0948 - accuracy: 0.9727 - va
Epoch 13/20
469/469 [=====] - 6s 14ms/step - loss: 0.0955 - accuracy: 0.9714 - va
Epoch 14/20
469/469 [=====] - 6s 13ms/step - loss: 0.0882 - accuracy: 0.9746 - va
Epoch 15/20
469/469 [=====] - 6s 13ms/step - loss: 0.0850 - accuracy: 0.9754 - va
Epoch 16/20
469/469 [=====] - 6s 14ms/step - loss: 0.0779 - accuracy: 0.9770 - va
Epoch 17/20
469/469 [=====] - 6s 14ms/step - loss: 0.0776 - accuracy: 0.9762 - va
Epoch 18/20
469/469 [=====] - 6s 13ms/step - loss: 0.0755 - accuracy: 0.9773 - va
Epoch 19/20
469/469 [=====] - 6s 13ms/step - loss: 0.0724 - accuracy: 0.9791 - va
Epoch 20/20
469/469 [=====] - 6s 13ms/step - loss: 0.0678 - accuracy: 0.9804 - va

```

```
In [ ]: score = model_3lbn.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

```


```

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```
Test score: 0.06718681752681732
```

```
Test accuracy: 0.9817000031471252
```



0.4 Five hidden layer

0.4.1 MLP + ReLU + ADAM + w/o Dropout + BN + 5Layer

```
In [ ]: initializer = tf.keras.initializers.he_normal(seed=None)
```

```
model_5lbn = Sequential()
```

```
model_5lbn.add(Dense(612, activation='relu', input_shape=(input_dim,), kernel_initializer=initializer))
model_5lbn.add(BatchNormalization())
```

```
model_5lbn.add(Dense(512, activation='relu', kernel_initializer=initializer) ) #layer 2
model_5lbn.add(BatchNormalization())
```

```
model_5lbn.add(Dense(312, activation='relu', kernel_initializer=initializer) ) #layer 3
model_5lbn.add(BatchNormalization())
```

```
model_5lbn.add(Dense(212, activation='relu', kernel_initializer=initializer) ) #layer 4
model_5lbn.add(BatchNormalization())
```

```
model_5lbn.add(Dense(15, activation='relu', kernel_initializer=initializer) ) #layer 5
model_5lbn.add(BatchNormalization())
```

```
model_5lbn.add(Dense(output_dim, activation='softmax'))
```

```
model_5lbn.summary()
```

```
Model: "sequential_10"
```

Layer (type)	Output Shape	Param #
dense_34 (Dense)	(None, 612)	480420
batch_normalization_25 (Batch Normalization)	(None, 612)	2448
dense_35 (Dense)	(None, 512)	313856
batch_normalization_26 (Batch Normalization)	(None, 512)	2048
dense_36 (Dense)	(None, 312)	160056
batch_normalization_27 (Batch Normalization)	(None, 312)	1248
dense_37 (Dense)	(None, 212)	66356
batch_normalization_28 (Batch Normalization)	(None, 212)	848
dense_38 (Dense)	(None, 15)	3195
batch_normalization_29 (Batch Normalization)	(None, 15)	60
dense_39 (Dense)	(None, 10)	160
Total params: 1,030,695		
Trainable params: 1,027,369		
Non-trainable params: 3,326		

```
In [ ]: model_5lbn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_5lbn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
Epoch 1/20
```

```
469/469 [=====] - 13s 27ms/step - loss: 0.2785 - accuracy: 0.9367 - val_loss: 0.1034 - val_accuracy: 0.9739
```

```
Epoch 2/20
```

```
469/469 [=====] - 12s 26ms/step - loss: 0.0934 - accuracy: 0.9739 - val_loss: 0.0626 - val_accuracy: 0.9814
```

```
Epoch 3/20
```

```
469/469 [=====] - 12s 27ms/step - loss: 0.0626 - accuracy: 0.9814 - val_loss: 0.0506 - val_accuracy: 0.9841
```

```
Epoch 4/20
```

```
469/469 [=====] - 12s 26ms/step - loss: 0.0506 - accuracy: 0.9841 - val_loss: 0.0484 - val_accuracy: 0.9841
```

```

Epoch 5/20
469/469 [=====] - 12s 26ms/step - loss: 0.0405 - accuracy: 0.9875 - va
Epoch 6/20
469/469 [=====] - 12s 26ms/step - loss: 0.0332 - accuracy: 0.9893 - va
Epoch 7/20
469/469 [=====] - 13s 28ms/step - loss: 0.0320 - accuracy: 0.9899 - va
Epoch 8/20
469/469 [=====] - 12s 26ms/step - loss: 0.0238 - accuracy: 0.9924 - va
Epoch 9/20
469/469 [=====] - 12s 26ms/step - loss: 0.0249 - accuracy: 0.9918 - va
Epoch 10/20
469/469 [=====] - 12s 26ms/step - loss: 0.0223 - accuracy: 0.9930 - va
Epoch 11/20
469/469 [=====] - 12s 26ms/step - loss: 0.0194 - accuracy: 0.9936 - va
Epoch 12/20
469/469 [=====] - 12s 26ms/step - loss: 0.0202 - accuracy: 0.9937 - va
Epoch 13/20
469/469 [=====] - 13s 27ms/step - loss: 0.0172 - accuracy: 0.9945 - va
Epoch 14/20
469/469 [=====] - 14s 29ms/step - loss: 0.0185 - accuracy: 0.9942 - va
Epoch 15/20
469/469 [=====] - 13s 28ms/step - loss: 0.0138 - accuracy: 0.9954 - va
Epoch 16/20
469/469 [=====] - 12s 26ms/step - loss: 0.0140 - accuracy: 0.9954 - va
Epoch 17/20
469/469 [=====] - 13s 27ms/step - loss: 0.0140 - accuracy: 0.9954 - va
Epoch 18/20
469/469 [=====] - 12s 27ms/step - loss: 0.0135 - accuracy: 0.9956 - va
Epoch 19/20
469/469 [=====] - 12s 27ms/step - loss: 0.0131 - accuracy: 0.9958 - va
Epoch 20/20
469/469 [=====] - 13s 28ms/step - loss: 0.0084 - accuracy: 0.9973 - va

```

```
In [ ]: score = model_5lbn.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))


```

```
vy = history.history['val_loss']
```

```
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0679909959435463

Test accuracy: 0.9822999835014343



0.4.2 MLP + ReLU + ADAM + BN + with Dropout + 5Layer

```
In [ ]: initializer = tf.keras.initializers.he_normal(seed=None)
```

```
model_5lbnd = Sequential()
```

```
model_5lbnd.add(Dense(612, activation='relu', input_shape=(input_dim,), kernel_initializer=initializer))
```

```
model_5lbnd.add(BatchNormalization())
```

```
model_5lbnd.add(Dropout(0.5))
```

```
model_5lbnd.add(Dense(512, activation='relu', kernel_initializer=initializer)) #layer 2
```

```
model_5lbnd.add(BatchNormalization())
```

```
model_5lbnd.add(Dropout(0.5))
```

```
model_5lbnd.add(Dense(312, activation='relu', kernel_initializer=initializer)) #layer 3
```

```
model_5lbnd.add(BatchNormalization())
```

```
model_5lbnd.add(Dropout(0.5))
```

```

model_5lbnd.add(Dense(212, activation='relu', kernel_initializer=initializer) ) #layer
model_5lbnd.add(BatchNormalization())
model_5lbnd.add(Dropout(0.5))

model_5lbnd.add(Dense(15, activation='relu', kernel_initializer=initializer) ) #layer
model_5lbnd.add(BatchNormalization())
model_5lbnd.add(Dropout(0.5))

model_5lbnd.add(Dense(output_dim, activation='softmax'))

model_5lbnd.summary()

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 612)	480420
batch_normalization_20 (Batch Normalization)	(None, 612)	2448
dropout_5 (Dropout)	(None, 612)	0
dense_29 (Dense)	(None, 512)	313856
batch_normalization_21 (Batch Normalization)	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0
dense_30 (Dense)	(None, 312)	160056
batch_normalization_22 (Batch Normalization)	(None, 312)	1248
dropout_7 (Dropout)	(None, 312)	0
dense_31 (Dense)	(None, 212)	66356
batch_normalization_23 (Batch Normalization)	(None, 212)	848
dropout_8 (Dropout)	(None, 212)	0
dense_32 (Dense)	(None, 15)	3195
batch_normalization_24 (Batch Normalization)	(None, 15)	60
dropout_9 (Dropout)	(None, 15)	0

```
dense_33 (Dense)                (None, 10)                160
=====
Total params: 1,030,695
Trainable params: 1,027,369
Non-trainable params: 3,326
-----
```

```
In [ ]: model_5lbnd.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_5lbnd.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
Epoch 1/20
469/469 [=====] - 14s 30ms/step - loss: 1.2392 - accuracy: 0.6089 - val_loss: 1.2392 - val_accuracy: 0.6089
Epoch 2/20
469/469 [=====] - 14s 29ms/step - loss: 0.5216 - accuracy: 0.8631 - val_loss: 0.5216 - val_accuracy: 0.8631
Epoch 3/20
469/469 [=====] - 14s 29ms/step - loss: 0.3929 - accuracy: 0.8996 - val_loss: 0.3929 - val_accuracy: 0.8996
Epoch 4/20
469/469 [=====] - 14s 29ms/step - loss: 0.3270 - accuracy: 0.9171 - val_loss: 0.3270 - val_accuracy: 0.9171
Epoch 5/20
469/469 [=====] - 13s 28ms/step - loss: 0.2912 - accuracy: 0.9244 - val_loss: 0.2912 - val_accuracy: 0.9244
Epoch 6/20
469/469 [=====] - 13s 29ms/step - loss: 0.2621 - accuracy: 0.9315 - val_loss: 0.2621 - val_accuracy: 0.9315
Epoch 7/20
469/469 [=====] - 14s 29ms/step - loss: 0.2509 - accuracy: 0.9339 - val_loss: 0.2509 - val_accuracy: 0.9339
Epoch 8/20
469/469 [=====] - 14s 29ms/step - loss: 0.2287 - accuracy: 0.9392 - val_loss: 0.2287 - val_accuracy: 0.9392
Epoch 9/20
469/469 [=====] - 13s 29ms/step - loss: 0.2177 - accuracy: 0.9425 - val_loss: 0.2177 - val_accuracy: 0.9425
Epoch 10/20
469/469 [=====] - 14s 29ms/step - loss: 0.2143 - accuracy: 0.9446 - val_loss: 0.2143 - val_accuracy: 0.9446
Epoch 11/20
469/469 [=====] - 14s 29ms/step - loss: 0.1945 - accuracy: 0.9482 - val_loss: 0.1945 - val_accuracy: 0.9482
Epoch 12/20
469/469 [=====] - 14s 29ms/step - loss: 0.1935 - accuracy: 0.9484 - val_loss: 0.1935 - val_accuracy: 0.9484
Epoch 13/20
469/469 [=====] - 13s 29ms/step - loss: 0.1842 - accuracy: 0.9506 - val_loss: 0.1842 - val_accuracy: 0.9506
Epoch 14/20
469/469 [=====] - 13s 29ms/step - loss: 0.1776 - accuracy: 0.9508 - val_loss: 0.1776 - val_accuracy: 0.9508
Epoch 15/20
469/469 [=====] - 14s 29ms/step - loss: 0.1720 - accuracy: 0.9523 - val_loss: 0.1720 - val_accuracy: 0.9523
Epoch 16/20
469/469 [=====] - 14s 29ms/step - loss: 0.1612 - accuracy: 0.9557 - val_loss: 0.1612 - val_accuracy: 0.9557
Epoch 17/20
469/469 [=====] - 13s 28ms/step - loss: 0.1627 - accuracy: 0.9561 - val_loss: 0.1627 - val_accuracy: 0.9561
Epoch 18/20
469/469 [=====] - 13s 28ms/step - loss: 0.1553 - accuracy: 0.9567 - val_loss: 0.1553 - val_accuracy: 0.9567
```

Epoch 19/20

469/469 [=====] - 13s 29ms/step - loss: 0.1518 - accuracy: 0.9589 - va

Epoch 20/20

469/469 [=====] - 14s 29ms/step - loss: 0.1494 - accuracy: 0.9589 - va

```
In [ ]: score = model_5lbnd.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])
        print('Test accuracy:', score[1])
```


```
fig,ax = plt.subplots(1,1)
ax.set_title('EpochsVS Loss')
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08210139721632004

Test accuracy: 0.9828000068664551



0.5 Conclusion

```
In [1]: from prettytable import PrettyTable
        x = PrettyTable()
        x.field_names = ["No. of layer", "Dropout", "Accuracy %"]

        x.add_row(["2", "NO", "98.03"])
        x.add_row(["2", "YES", "98.40"])
        x.add_row(["3", "NO", "97.76"])
        x.add_row(["3", "YES", "98.17"])
        x.add_row(["5", "NO", "98.22"])
        x.add_row(["5", "YES", "98.28"])

        print(x)
```

No. of layer	Dropout	Accuracy %
2	NO	98.03
2	YES	98.40
3	NO	97.76
3	YES	98.17
5	NO	98.22
5	YES	98.28

1. Observed that using dropouts there is slight increase in the accuracy
2. 5 hidden layer model has given 98% accuracy and no much increment with/without dropout as dataset is small